


INC-2343 datadog-prod-in-awsstprod-reached-open-files-limit

 Eric Swanson

Limit is 2,000 parallel tailed logs per agent.

- Identifying highest instant volume sources will be helpful in reducing the immediate impact
 - **NOTE:** Adjust focus to number of FILES across all agents sharing that open file limit, NOT the individual logfile message volume.
- removing lower volume errors can free up headroom and also do cleanup we should be getting to for hygiene anyway.
 - 240 errors per hour vs 411k logs in 10 minutes.....
- <https://docs.datadoghq.com/agent/logs/?tab=tailfiles>
 - The maximum number of files that can be tailed in parallel.
- <https://github.com/dbt-labs/helm-charts/blob/0a4284ca62da696dcc77ce920e99f92d2194791a/charts/datadog/values.yaml#L40-L42>


```
env :  
  - name : DD_LOGS_CONFIG_OPEN_FILES_LIMIT  
    value : '2000'
```

- pod_name: datadog-agent-x8sr9
- asg: usaa-prod-run-orc-worker-group20240912193839671000000004


 [Orchestration Overview](#)

usaa-prod orc-dispatch Log count by kube_node and index

2d Apr 12 - Apr 14

No Data


Logfile Message Count Spikes by Filename in usaa-prod



usaa-prod logs excluding high volume spike message contents

2d

Apr 12 - Apr 14



Excluded Logs

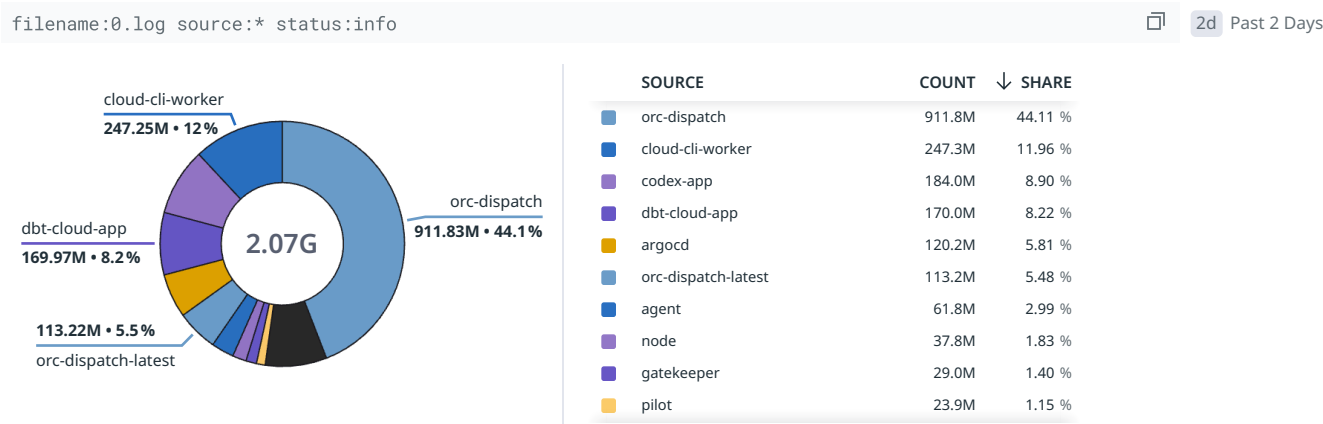
2d

Apr 12 - Apr 14



COUNT	-VOLUME	MESSAGE
<div><div></div><div>≈363K</div></div>	<div><div></div><div>98K</div></div>	HH:mm:ss target * is ::
<div><div></div><div>≈127K</div></div>	<div><div></div><div>31.2K</div></div>	HH:mm:ss set * session variables complete
<div><div></div><div>≈126K</div></div>	<div><div></div><div>34.4K</div></div>	HH:mm:ss target * is :*:
<div><div></div><div>≈118K</div></div>	<div><div></div><div>29.6K</div></div>	HH:mm:ss EXCEPTION: Model .. does not exist
<div><div></div><div>≈74.8K</div></div>	<div><div></div><div>18.7K</div></div>	HH:mm:ss stage load logging complete
<div><div></div><div>≈67.2K</div></div>	<div><div></div><div>18.5K</div></div>	HH:mm:ss EXCEPTION: Cant extract ETL audit columns ETL_CREA_NR, ETL_CREA_GMT_TS
<div><div></div><div><1K</div></div>	<div><div></div><div><1K</div></div>	HH:mm:ss target relation is : "PENTCAOITDW"."*"."*":

0.log info volume by top 10 services



► Paging Alert Events

Open File Count Investigation

agent log ingestion patterns


1h 1:00 am - 2:00 am

COUNT	~VOLUME	STATUS	SERVICE	MESSAGE	
-------	---------	--------	---------	---------	--




No matching results found


CORE | INFO | (pkg/logs/schedulers/ad/scheduler.go:73 in Schedule) | Received a new logs config: container_collect_all 1h 1:00 am - 2:00 am




pkg/logs/launchers/file/launcher.go starting a new tailer

1h 1:00 am - 2:00 am






CPU usage (%)




USAA Agent pod CPU Requests, Limits and Usage during 2 AM Paging alert for Max Open Files limit reached 3h12:00 am - 3:00 am




Disk I/O

3h 12:00 am - 3:00 am




Network Rate

3h 12:00 am - 3:00 am




Pod Ephemeral Storage Usage

3h 12:00 am - 3:00 am



USAA Agent pod open file count during 2AM Paging Alert window

3h 12:00 am - 3:00 am



status:info dbt_instance:usaa-prod service:orc-dispatch @filename:client.py

10s 2:14 am - 2:14 am

COUNT	~VOLUME	@FILENAME	MESSAGE
-------	---------	-----------	---------



No matching results found

Additional Details


Impact window: Agent check status by check name

1h 12:58 am - 1:58 am

No Data


dbt-cloud-app account feature log volume around impact window

2h 12:51 am - 2:58 am




Top 5 orc-dispatch log source filenames by volume

2h 12:51 am - 2:58 am




USAA Prod Datadog Agent max open file count 4h rollup past 4w

28d Past 28 Days



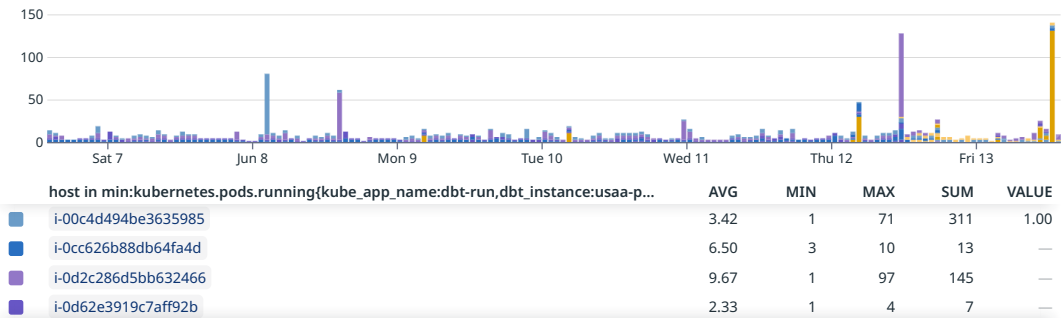
15 minute sum of minimum kubernetes.pods.running (Gauge) by host for dbt-run in usaa-prod

1d Apr 13 - Apr 14



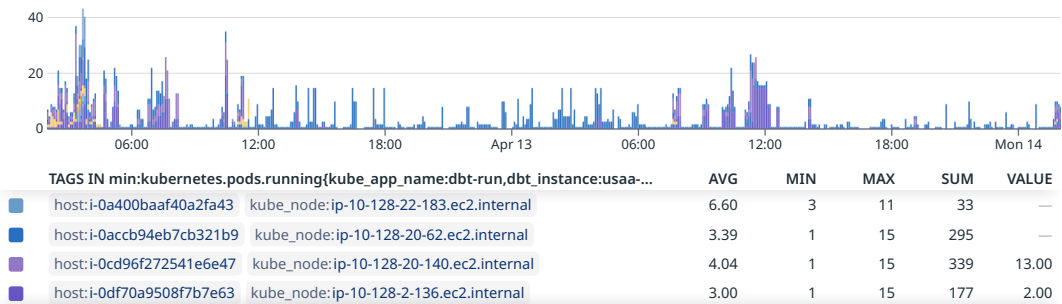
15 minute sum of minimum kubernetes.pods.running (Gauge) by host for dbt-run in usaa-prod

7d Past 7 Days




15 minute sum of minimum kubernetes.pods.running (Gauge) by host for dbt-run in usaa-prod

2d Apr 12 - Apr 14




15 minute sum of minimum kubernetes.pods.running (Gauge) by host and kube_node (ip) for dbt-run in usaa-prod 1d Past 1 Day




Hourly count of st TRUNCATED orc-dispatch Error Logs by dbt_instance

5d Past 5 Days




orc-dispatch Host Warning Logs by Container


1h rollup of warning log count by container name for host with most log volume



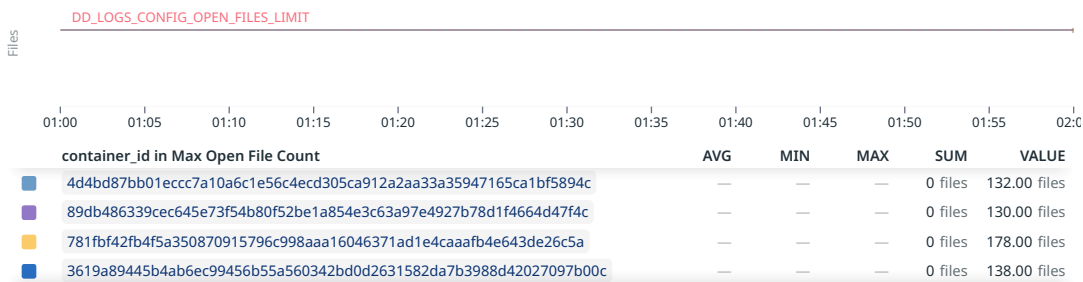
service:orc-dispatch* filename:0.log container_name:run \$dbt_instance status:info, service:orc-dispatch* filename:0.log container_name:run




USAA Prod Datadog Agent max open file 1h rollup count 1w



USAA Prod Datadog Agent max open file 2h rollup count 2w by agent id



CPU Usage by Pod



1h 9:22 am - 10:23 am

MESSAGE

2025-05-08 05:54:44 UTC | CORE | WARN | (pkg/logs/launchers/file/provider/file_provider.go:232 in FilesToTail) | Reached the limit on the maximum number of files in use: 500

TAGS

HOST

i-020476c2d34bc2e88

SERVICE

agent

SOURCE

agent

CONTAINER NAME

agent

DOCKER IMAGE

gcr.io/datadoghq/agent

POD NAME

datadog-agent-g8b2w


availability-zone:us-east-1c aws_account:850607893674 aws_autoscaling_groupname:usaa-prod-run-orc-worker-group202409121938396710000... +51

EVENT ATTRIBUTES

Show More

Count by dbt_instance over "source:datadog dbt_instance:usaa* @evt.type:"Agent Startup""

14d Past 14 Days



Count by dbt_instance over "source:datadog dbt_env:prod @evt.type:"Agent Startup""

1d Past 1 Day

